# Assignment 3, 2019

Released: 2 May.
Submit test data (individual): Monday 20 May at 23:00
Submit compiler (team effort): Wednesday 29 May at 23:00

## Objectives

To build a better understanding of a compiler's back-end, code generation, symbol tables, run-time structures, and optimization. To practice cooperative, staged software development.

## Background and context

The task is to write a compiler for a procedural language, Goat. The compiler translates source programs to the assembly language of a target machine Oz[1]. These programs can then be run on a provided Oz emulator.

In an earlier stage, you wrote a parser for Goat. You may choose to start from that parser, or alternatively, start from one that has been (or will be) made available. In either case, correctness of the compiler, including the parser, is your responsibility.

This final stage also involves the completion of semantic analysis, code generation, and further optional tasks. The implementation language must be Haskell.

## The source language: Goat

The syntax of Goat is already known to you from Stage 1. A Goat program lives in a single file and consists of a number of procedure definitions. There are no global variables. One (parameter-less) procedure must be named "main". Procedure parameters can be passed by value or by reference.

There are three *base types*, namely *int*, *float*, and *bool*. The first two are *numeric* types and allow for arithmetic and comparison operators. Goat also has static arrays and matrices (2-dimensional arrays). The "write" command can print integers, floating point numbers, Booleans and, additionally, strings.

### Syntax

The following are reserved words: `begin`, `bool`, `call`, `do`, `else`, `end`, `false`, `fi`, `float`, `if`, `int`, `od`, `proc`, `read`, `ref`, `then`, `true`, `val`, `while`, `write`. The lexical rules were given in the specification for Stage 1.

An int literal is a non-empty sequence of digits. A float literal is a sequence of one or more digits, followed by a decimal point, and another sequence of one or more digits. A Boolean constant is `false` or `true`. A string constant (as can be used by "`write`") is a sequence of characters between double quotes. The sequence itself cannot contain double quotes or escape

---

[1]Not to be confused with the multi-paradigm language Oz, see `http://en.wikipedia.org/wiki/Oz_(programming_language)`.

characters, except it may include "\n", representing a newline character. That is, a backslash not preceding an 'n' is just considered part of the string.

The arithmetic binary operators associate to the left, and unary operators have higher precedence (so for example, '-5+6' and '4-2-1' both evaluate to 1).

A Goat program consists of one or more procedure definitions. Each definition consists of

1. the keyword `proc`,
2. a procedure header,
3. a procedure body,
4. the keyword `end`,

in that order.

The header has two components (in this order):

1. An identifier—the procedure's name.
2. A (possibly empty) comma-separated list of formal parameters within a pair of parentheses (so the parentheses are always present).

Each formal parameter has three components:

1. A parameter passing indicator (`val` or `ref`).
2. A type (`bool`, `float` or `int`).
3. An identifier.

The procedure body consists of zero or more local variable declarations, followed by the keyword `begin`, followed by a non-empty sequence of statements. A variable declaration consists of one of the keywords `bool`, `float` or `int`, followed by an identifier, terminated with a semicolon. Before the semicolon may be a shape indicator, which is either `[n]` or `[m,n]`, where `m` and `n` are integer literals (the former case indicates the declaration of an array of length `n`, the latter an `m` by `n` matrix). There may be any number of variable declarations, given in any order.

An atomic statement has one of the following forms:

```
<id> := <expr> ;
<id> [ <expr> ] := <expr> ;
<id> [ <expr> , <expr> ] := <expr> ;
read <id> ;
read <id> [ <expr> ] ;
read <id> [ <expr> , <expr> ] ;
write <expr> ;
call <id> ( <expr-list> ) ;
```

where `<expr-list>` is a (possibly empty) comma-separated list of expressions.

A composite statement has one of the following forms:

```
if <expr> then <stmt-list> fi
if <expr> then <stmt-list> else <stmt-list> fi
while <expr> do <stmt-list> od
```

where `<stmt-list>` is a non-empty sequence of statements, atomic or composite. (Note that semicolons are used to *terminate* atomic statements, so that a sequence of statements—atomic or composite—does not require any punctuation to *separate* the components.)

An expression has one of the following forms:

```
<id>
<id> [ <expr> ]
<id> [ <expr> , <expr> ]
<const>
( <expr> )
<expr> <binop> <expr>
<unop> <expr>
```

The list of operators is

```
||
&&
!
= != < <= > >=
+ -
* /
-
```

Here `!` is a unary prefix operator (Boolean negation), and the bottom "`-`" is a unary prefix operator (unary minus). All other operators are binary and infix. All the operators on the same line have the same precedence, and the ones on later lines have higher precedence; the highest precedence being given to unary minus. The six relational operators are non-associative (so for example, `x = y = z` is not well-formed). The six remaining binary operators are all left-associative. The relational operators yield Boolean values `true` or `false`, according as the relation is true or not. The Boolean operators `||` and `&&` are for disjunction ('or') and conjunction ('and'), respectively.

The language supports comments, which start at a # character and continue to the end of the line. White space is not significant.

## Static semantics

We define a *scalar* to be an object that contains a single value. This value can be inspected and updated. Variables, array elements $id[e]$, and matrix elements $id[e_1, e_2]$ are scalars. Goat does not allow an array *as a whole* to be assigned or passed as a parameter.[2] Only *selective updating* is possible. In other words, if $id$ has been declared as an array then $id$ is illegal as an expression on its own. It follows that every valid expression has type `bool`, `int`, or `float`.

All defined procedures must have distinct names. A defined procedure does not have to be called anywhere, and the definition of a procedure does not have to precede the (textually) first call to the procedure. Procedures can use (mutual) recursion. For each procedure, the number of actual parameters in a call must be equal to the number of formal parameters in the procedure's definition.

A formal procedure parameter is treated as a local variable. The scope of a declared variable (or of a formal procedure parameter) is the enclosing procedure definition. A variable must be declared (exactly once) before used. Similarly a list of formal parameters must all be distinct. However, the same variable/parameter name can be used in different procedures. An expression in an argument/actual parameter position where the parameter passing method is "by reference" must be a scalar (variable, array element, or matrix element).

In an array $(id[n])$ or matrix $(id[m, n])$ declaration, the sizes $m$ and $n$ must be positive integers. Each use of an array variable must include exactly one index expression, and each use of a matrix variable must include exactly two. Array and matrix identifiers obey the same scope

---

[2]We leave this as an optional extension task.

rules as other variables. An array or matrix cannot be passed as a whole entity to a procedure, but an array or matrix element, such as `a[3,2]` can be used as an actual parameter, passed by value or reference. Since the array bounds are known at declaration-time, semantic analysis *could* do limited bounds checking at compile-time, but there is no requirement to do this.

The language is statically typed, that is, each variable and parameter has a fixed type, chosen by the programmer. The type rules for expressions are as follows:

- The type of a Boolean constant is `bool`.

- The type of an integer constant is `int`.

- The type of a float constant is `float`.

- The type of an expression $id$ is the variable $id$'s declared type ($id$ cannot be an array or matrix variable).

- For an array element $id[e]$, $e$ must have type `int` and the type of the whole expression is the type given in the declaration of $id$.

- For a matrix element $id[e_1, e_2]$, $e_1$ and $e_2$ must have type `int`, and the type of the whole expression is the type given in the declaration of $id$.

- Arguments of the logical operators must be of type `bool`. The result of applying these operators is of type `bool`.

- The two operands of `=` must have the same (base) type; the same goes for `!=`. The result is Boolean.

- The two operands of any other comparison operator must either have the same (base) type, or one must be `int` and the other `float` (in which case the compiler must then convert the integer to a float before the comparison can happen). The ordering on Boolean values is defined by $x \leq y$ iff $x$ is `false` or $y$ is `true` (or both these hold). As usual, $x < y$ iff $x \leq y \wedge x \neq y$. The result of a comparison is Boolean, irrespective of the operand types.

- The two operands of a binary arithmetic operator must either have the same numeric type, or one must be `int` and the other `float` (in which case the compiler must convert the integer to a float before the comparison). If both operands have type `int`, the result type is `int`; otherwise it is `float`.

- The operand of unary minus must be of type `int` or `float`, and the result type is the same as the operand's type.

The type rules for statements are as follows:

- In assignment statements, a variable, array element, or matrix element on the left-hand side must have the same type as the expression on the right-hand side, with one exception: an `int` expression may be assigned to a `float` variable (the compiler must convert the value to be assigned from `int` to `float`). Note that arrays and matrices can only be updated selectively, that is, only variables, array *elements*, and matrix *elements* can be assigned to.

- Conditions in `if` and `while` statements must be of type `bool`. Their bodies must be well-typed sequences of statements.

- In procedure calls, the type of an actual parameter must be the type of the corresponding formal parameter, or (for passing "by value"), the actual parameter may be of type `int` and the formal parameter of type `float` (in which case the compiler must convert the integer to a float).

- `read` takes a scalar (variable, array element, or matrix element).

- The argument to `write` must be a well-typed expression *or* a string literal.

Every procedure in a program must be well-typed, that is, its body must be a sequence of well-typed statements. Every program must contain a procedure of arity 0 named "main".

### Dynamic semantics

Numerical variables are automatically initialised to 0, and Boolean variables to `false`. This extends to arrays and matrices. The evaluation of an expression $e_1/e_2$ results in a runtime error if $e_2$ evaluates to 0.

The binary logical operators are *non-strict* and their arguments are evaluated from left to right. For example, `5 < 8 || 5 > 8/0` yields `true` rather than a runtime error.

If a Goat program reads or writes outside the bounds of an array or matrix, the behaviour is undefined. (There is no requirement that out-of-bounds indices are detected at runtime, though this may be implemented as an optional extension, as discussed below).

`write` prints `int`, `float` and `bool` expressions to `stdout` in their standard syntactic forms, with no additional whitespace or newlines. If `write` is given a string literal, it prints out the characters of the string to `stdout`, with `\n` resulting in a newline character being printed. Similarly, `read` reads a `int`, `float` and `bool` literal from `stdin` and assigns it to a variable, array element or matrix element. If the user input is not valid, execution terminates.[3]

The procedure "main" is the entry point, that is, execution of a program comes down to execution of a call to "main".

The language allows for two ways of passing parameters. Call by value (`val`) is a copying mechanism. For each parameter $e$ passed by value, the called procedure considers the corresponding formal parameter $v$ a local variable and initialises this local variable to $e$'s value.

Call by reference (`ref`) does not involve copying. Instead the called procedure is provided with the *address* of the actual parameter (which must be a variable $z$, array element $a[e]$ or matrix element $m[e_1, e_2]$), and the formal parameter $v$ is considered a synonym for the actual parameter.[4]

Some subtleties of parameter passing come about as the result of *aliasing*. Consider the program on the right. If `<method>` is `val`, the program will print '4'. If it is `ref`, the program will print '8'. If both arguments were passed by value, the result would have been '3'.

```
proc main()
int z;
    z := 3;
    p(z,z);
    write z;
end

proc p(ref int x, <method> int y)
    x := 4;
    y := y + x;
end
```

---

[3]Various instructions in the Oz assembly language can be used to take care of these rules for you.

[4]In terms of stack slots in the frame allocated for a procedure call, one slot is needed for a parameter passed by value. The parameter is simply treated as a local variable. A parameter passed by reference also needs one slot, but in this case, the *address* of the parameter is what is stored, and indirect addressing must be used to access the variable.
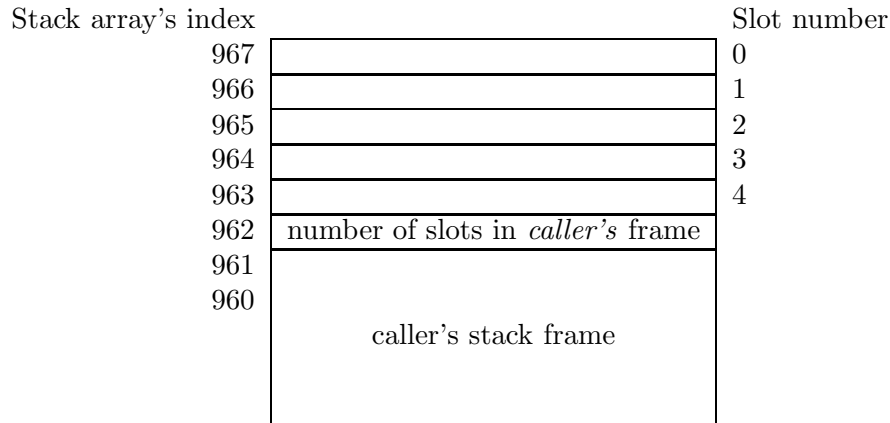
Figure 1: A stack frame on the 'stack' array

The rest of the semantics should be obvious—it follows standard conventions. For example, the execution of `while` $e$ `do` $ss$ `od` can be described as follows. First evaluate $e$. If $e$ evaluates to `false`, the statement is equivalent to a no-op (a statement that does nothing). Otherwise the statement is equivalent to '$ss$ `while` $e$ `do` $ss$ `od`'.

## The target language: Oz

Oz is an artificial target machine that closely resembles intermediate representations used by many compilers. An emulator for the Oz machine is supplied to you as a C program. This emulator reads Oz source files (which should be the output of your compiler) and executes them without further compilation. You are not required to modify the emulator, or understand its inner workings, and may treat it as a "black box" (although you may want to study the source code for your own benefit).

Oz has 1024 registers named r0, r1, r2, ... r1023. This is effectively an unlimited set of registers, and your compiler may treat it as such; your compiler may generate register numbers without checking whether they exceed 1023. Every register may contain a value of type `int` or `float` (referred to as `real` in the emulator).

Oz also has an area of main memory representing the stack, which contains zero or more stack frames. Each stack frame contains a number of stack slots. Each stack slot may contain a value of any type (it also holds type information about the value, for validation purposes), and you specify a stack slot by its stack slot number.

Figure 1 shows the "top" of the stack at some point. A stack pointer keeps track of the highest index used in the array, and stack slots can be accessed relative to this. Notice how the numbering of slots runs in the opposite direction of the array indices. In the example, the current stack frame (or activation record) has five slots (actually, it has one additional slot, used for remembering the size of the stack frame that will resume as current, once the active procedure returns).

An Oz program consists of a sequence of instructions, some of which may have labels. Although the emulator does not require it, good style dictates that each instruction should be written on its own line. As in most assembly languages, you can attach a label to an instruction by preceding it with an identifier (the name of the label) and a colon. The label and the instruction may be on the same line or different lines. Again, good style suggests that labels are somehow made to stand out in the generated Oz code. Identifiers have the same format in Oz as in Goat.

6

The following lists the relevant opcodes of Oz (there are a few more), and for each opcode, shows how many operands it has, and what they are. *The destination operand is always the leftmost operand.*

```
push_stack_frame   framesize
pop_stack_frame    framesize
                                        #  C analogues:
store              slotnum, rI          #       x =  rI
load               rI, slotnum          #      rI =   x
load_address       rI, slotnum          #      rI =  &x
load_indirect      rI, rJ               #      rI = *rJ
store_indirect     rI, rJ               #     *rI =  rJ

int_const          rI, intconst
real_const         rI, realconst
string_const       rI, stringconst

add_int            rI, rJ, rK           #  rI =  rJ + rK
add_real           rI, rJ, rK           #  rI =  rJ + rK
add_offset         rI, rJ, rK           #  rI =  rJ + rK
sub_int            rI, rJ, rK           #  rI =  rJ - rK
sub_real           rI, rJ, rK           #  rI =  rJ - rK
sub_offset         rI, rJ, rK           #  rI =  rJ - rK
mul_int            rI, rJ, rK           #  rI =  rJ * rK
mul_real           rI, rJ, rK           #  rI =  rJ * rK
div_int            rI, rJ, rK           #  rI =  rJ / rK
div_real           rI, rJ, rK           #  rI =  rJ / rK
neg_int            rI, rJ               #  rI = -rJ
neg_real           rI, rJ               #  rI = -rJ

cmp_eq_int         rI, rJ, rK           #  rI = rJ == rK
cmp_ne_int         rI, rJ, rK           #  rI = rJ != rK
cmp_gt_int         rI, rJ, rK           #  rI = rJ >  rK
cmp_ge_int         rI, rJ, rK           #  rI = rJ >= rK
cmp_lt_int         rI, rJ, rK           #  rI = rJ <  rK
cmp_le_int         rI, rJ, rK           #  rI = rJ <= rK
cmp_eq_real        rI, rJ, rK           #  rI = rJ == rK
cmp_ne_real        rI, rJ, rK           #  rI = rJ != rK
cmp_gt_real        rI, rJ, rK           #  rI = rJ >  rK
cmp_ge_real        rI, rJ, rK           #  rI = rJ >= rK
cmp_lt_real        rI, rJ, rK           #  rI = rJ <  rK
cmp_le_real        rI, rJ, rK           #  rI = rJ <= rK

and                rI, rJ, rK           #  rI = rJ && rK
or                 rI, rJ, rK           #  rI = rJ || rK
not                rI, rJ               #  rI = !rJ

int_to_real        rI, rJ               #  rI = (float) rJ
move               rI, rJ               #  rI = rJ
```

```
        branch_on_true    rI, label              #  if (rI) goto label
        branch_on_false   rI, label              #  if (!rI) goto label
        branch_uncond     label                  #  goto label


        call              label
        call_builtin      builtin_function_name
        return
        halt


        debug_reg         rI
        debug_slot        slotnum
        debug_stack
```

The `push_stack_frame` instruction creates a new stack frame. Its argument is an integer specifying how many slots the stack frame has; for example the instruction `stack_frame 5` creates a stack frame with five slots numbered 0 through 4. (In the emulator, it also reserves an extra slot, slot 5, to hold the size of the previous stack frame, for error detection purposes.)

The `pop_stack_frame` instruction deletes the current stack frame. Its argument is an integer specifying how many slots that stack frame has; it must match the argument of the `push_stack_frame` instruction that created the stack frame being popped.

The `store` instruction copies a value from the named register to the stack slot with the given number. The `load` instruction copies a value from the stack slot with the given number to the named register. When a slot is first created, Oz marks it as uninitialized. An attempt to load an uninitialized value results in a runtime error.

The `load_address` instruction can be used by a caller to facilitate call by reference. The called procedure, having stored the address in the current stack frame, can then access and change the content of that address, by moving the address to a register and using `load_indirect` and `store_indirect`.

The `add_offset` and `sub_offset` instructions calculate addresses based on the offset from a given stack slot. They are useful when array or matrix elements need to be accessed or updated. The instruction '`add_offset rI rJ rK`' assumes that `rJ` holds an address, and `rK` holds an integer offset to be added to that address, the result being placed in `rI` (and similarly for `sub_offset`). Note that the Oz emulator is designed so that slot numbers grow in the opposite direction to how addresses grow, so `sub_offset` is appropriate when you want to *add* offsets to slot numbers, see Figure 1.

The `int_const`, `real_const`, and `string_const` instructions all load a constant of the specified type to the named register. The format of the constants is exactly the same as in Goat.

The `add_int`, `add_real`, `sub_int`, `sub_real`, `mul_int`, `mul_real`, `div_int`, `div_real`, `neg_int`, and `neg_real` instructions perform arithmetic. The first part of the instruction name specifies the arithmetic operation, while the second part specifies the shared type of all the operands.

The `cmp_eq_int`, `cmp_ne_int`, `cmp_gt_int`, `cmp_ge_int`, `cmp_lt_int` and `cmp_le_int` instructions, and their equivalents for floats, perform comparisons, generating integer results. The middle part of the instruction name specifies the comparison operation, while the last part specifies the shared type of both input operands.

The `and`, `or` and `not` instructions each perform the "logical" operation of the same name. Oz represents boolean values as integers, taking 0 to represent 'false' and a non-zero integer to represent 'true'.

The `int_to_real` instruction converts the integer in the source register to a floating point number in the destination register.

The `move` instruction copies the value in the source register (which may be of any type) to the destination register.

The `branch_on_true` instruction transfers control to the specified label if the named register contains a non-zero integer value. The `branch_on_false` instruction transfers control to the specified label if the named register contains 0. The `branch_uncond` instruction always transfers control to the specified label.

The `call` instruction calls the procedure whose code starts with the label whose name is the operand of the instruction, while the `call_builtin` instruction calls the built-in function whose name is the operand of the instruction. Procedures and functions take their first argument from register r0, their second from r1, and so on. During the call, the procedure may destroy the values in all the registers, so they contain nothing meaningful when the procedure returns. The exception is that the built-in functions that return a value, such as the read functions, put their return value in r0 (see the example in Figure 3). When the called procedure executes the `return` instruction, execution continues with the instruction following the call instruction.

The following are all built-in functions: `read_int`, `read_real`, `read_bool`, `print_int`, `print_real`, `print_bool`, and `print_string`. (There are a few other built-ins that you will not need.) The read functions take no argument. They read a value of the indicated type (using C's scanf function) from standard input, and return it in r0. The function `read_bool` accepts the strings "true" and "false" and returns a 1 and a 0, respectively. The print functions take a single argument of the named type in r0, and print it to standard output; they return nothing.

Each `call` instruction pushes the return address (the address of the instruction following it) onto the stack. The `return` instruction transfers control to the address it pops off the stack.

The `halt` instruction stops the program.

Oz also supports comments, which start at a # character and continue until the end of the line. It can be useful to have the code generator insert comments, as in the example below.

The `debug_reg`, `debug_slot` and `debug_stack` instructions are Oz's equivalent of debugging printfs in C programs: they print the value in the named register or stack slot or the entire stack. They are intended for debugging only; your submitted compiler should not generate them. If your code generator generates Oz code that does the wrong thing and you cannot sort out why, you can manually insert these instructions to better see what goes wrong. Calling the emulator with an '-i' option gives a trace of execution.

Figure 2 shows the source program `gcd.gt`, and Figure 3 shows one possible translation. The Oz emulator starts execution with the first instruction in the program and stops when it executes the halt instruction. Note that the generated code therefore starts with a fixed two-instruction sequence that represents the Oz runtime system: the first instruction calls main, while the second (executed when main returns) is a halt instruction.

## Summary of Tasks, Suggestions

The compiler should take the name of a source file on the command line. It should write the corresponding target program to standard output, or report errors. If an error has been discovered in any phase of compilation, no (partial) code should be written. The executable compiler must be called `Goat`. On success, it should use an exit code of 0. Upon failure, it should use some non-0 exit code.

You already have a working parser, and if not, you can use the supplied one (but note that, in any case, correctness of the parser is your responsibility). There is no requirement to submit the pretty-printer, so it does not matter if you have to make changes to the AST that invalidate your pretty-printer. If you want to emit helpful comments in the Oz code, as shown in Figure 3, parts of your pretty-printer may come in handy.

```
proc main()
    int x;
    int y;
    int temp;
    int quotient;
    int remainder;

    write "Input two positive integers: ";

    read x;
    read y;

    write "\n";

    if x < y then
        temp := x;
        x := y;
        y := temp;
    fi

    write "The gcd of ";
    write x;
    write " and ";
    write y;
    write " is ";

    quotient := x / y;
    remainder := x - quotient * y;

    while remainder > 0 do
        x := y;
        y := remainder;
        quotient := x / y;
        remainder := x - quotient * y;
    od

    write y;
    write "\n";
end
```

Figure 2: The Goat program gcd.gt

The semantic analysis phase consists of a lot of checking that well-formedness conditions are met. The code generation phase consists of generating correct Oz code from the AST. Of these, well-formedness checking is arguably the part that has the lowest learning-outcome benefit for the time invested. The marking scheme encourages you to concentrate on code generation, and then deal with the correct handling of ill-formed programs as time allows.

You are encouraged to work stepwise and increase the part of the language you can handle as you go. It makes sense to write a module SymTable that offers the symbol table services. The

```
        call proc_main
        halt
proc_main:
  # prologue
        push_stack_frame 5
        int_const r0, 0
  # initialise int val quotient
        store 0, r0
  # initialise int val remainder
        store 1, r0
  # initialise int val temp
        store 2, r0
  # initialise int val x
        store 3, r0
  # initialise int val y
        store 4, r0
  # write string
        string_const r0, "Input two positive integers: "
        call_builtin print_string
  # read x
        call_builtin read_int
        store 3, r0
  # read y
        call_builtin read_int
        store 4, r0
  # write string
        string_const r0, "\n"
        call_builtin print_string
  # if x < y
        load r0, 3
        load r1, 4
        cmp_lt_int r0, r0, r1
        branch_on_true r0, label_0
        branch_uncond label_1
label_0:
  # temp := x
        load r0, 3
        store 2, r0
  # x := y
        load r0, 4
        store 3, r0
  # y := temp
        load r0, 2
        store 4, r0
label_1:
  # write string
        string_const r0, "The gcd of "
        call_builtin print_string
  # write x
        load r0, 3
        call_builtin print_int
  # write string
        string_const r0, " and "
        call_builtin print_string
  # write y
```

Figure 3: Translated program (first part)

```
    load r0, 4
    call_builtin print_int
  # write string
    string_const r0, " is "
    call_builtin print_string
  # quotient := x / y
    load r0, 3
    load r1, 4
    div_int r0, r0, r1
    store 0, r0
  # remainder := x - (quotient * y)
    load r0, 3
    load r1, 0
    load r2, 4
    mul_int r1, r1, r2
    sub_int r0, r0, r1
    store 1, r0
  # while remainder > 0
label_2:
    load r0, 1
    int_const r1, 0
    cmp_gt_int r0, r0, r1
    branch_on_true r0, label_3
    branch_uncond label_4
label_3:
  # x := y
    load r0, 4
    store 3, r0
  # y := remainder
    load r0, 1
    store 4, r0
  # quotient := x / y
    load r0, 3
    load r1, 4
    div_int r0, r0, r1
    store 0, r0
  # remainder := x - (quotient * y)
    load r0, 3
    load r1, 0
    load r2, 4
    mul_int r1, r1, r2
    sub_int r0, r0, r1
    store 1, r0
    branch_uncond label_2
label_4:
  # write y
    load r0, 4
    call_builtin print_int
  # write string
    string_const r0, "\n"
    call_builtin print_string
  # epilogue
    pop_stack_frame 5
    return
```

Figure 4: Translated program (remaining part)

`Data.Map` and `Data.Map.Strict` libraries offer a dictionary type that may be useful. A module `Analyze` could do the semantic analysis of the AST, and it might want to store information in the symbol table. Consider carefully whether you want to let the semantic analysis phase do all of the static analysis, or whether it should leave the type checking to the code generator. On the one hand, it is nice to have the code generator work on what it knows to be a well-typed program. On the other hand, the code generator needs to reconstruct much of the type inference anyway, because it needs to know when to insert type conversion code.

Work on getting the AST ready for code generation quickly—you can always add the well-formedness checks later, as time permits. A module `Codegen` can be responsible for code generation. It will also want to interact with the symbol table.

A possible approach to implementing `Goat` incrementally is as follows (note that some static analysis is needed from the outset—types need to be determined for code generation):

1. Get the compiler working for the subset that consists of expressions (not including arrays and matrices) and the `write` statement. Assume that procedures do not take any arguments and cannot use recursion (no procedure calls), so that `main` works. Now you should be able to compile `hello.gt` from the Kid examples given in Stage 1.

2. Add the `read` statement, and assignments. Now you should be able to compile `io.gt` and `asg.gt` from the same repository.

3. Add compound statements (`if` and `while`). Now you should be able to compile `gcd.gt`.

4. Add procedure arguments and procedure calls, but initially for pass-by-value only.

5. Add reference parameters.

6. Add arrays and matrices.

7. Complete static semantic analysis.

If you want to extend the task, here are some ideas:

- (Not much of an extension, really.) Add run-time discovery of bounds violations for arrays and matrices.

- Add optimisations based on peephole analysis (even a simple version can be quite effective). The emulator will provide statistics if called with an '-s' option.

- On top of the previous extension, for a more substantial challenge, add a static analysis to (1) remove unnecessary array and matrix bounds checks and (2) find definite bounds violations already at compile-time. The simplest solution would be to perform constant propagation. A more sophisticated approach would be an interval analysis that tracks which values integer expressions can take, over-approximating a set of values as an interval.

- Extend the source language so that arrays and matrices can be passed to procedures. Make appropriate syntactic and semantic design decisions.

- Add C-style functions to the source language.

The following directories are, or will be made, available on the LMS:

- `oz/` contains the Oz emulator. The make file will generate an executable called `oz`.

- `parser/` contains a `Goat` parser which is believed to be correct.

- `simple_tests/` contains a number of small Goat programs for testing.

- `contributed_tests/` will later contain Goat programs submitted by you. (Be prepared for the possibility that some of these may not actually be valid Goat, or they may not behave the way their authors assumed.)

## Procedure and assessment

The project may be solved in the teams, continuing from Stage 1. Each team should only submit once (under one of the members' name). If your team has changed since Stage 1, or if you are interested in taking on an extra team member, please let me know.

**By 20 May,** submit a single Goat program, which will be entered into a collection of test cases that will be made available to all. The program should be (syntactically, type, etc.) correct, but its runtime behaviour does not matter (whether it terminates, asks for input, divides by zero or whatever). Call your program `user.gt`, where `user` is your unimelb userid, and submit a separate file `user.in` with the intended input to `user.gt`, if it requires input. For this test submission stage, use `submit COMP90045 3a` to submit.

**By 29 May,** submit the code. There should be a `Makefile`, so that a `make` command generates `Goat`. Do not submit any files that are generated by a program generator such as alex or happy. Instead your `Makefile` should generate those files from alex or happy specifications that you submit. Also, do not submit `oz.c` or other files related to Oz. For this last stage, use `submit COMP90045 3b` to submit. It is possible to submit late, using `submit COMP90045 3b.late`, but late submissions will attract a penalty of 2 marks per calendar day late.

This project counts for 14 of the 30 marks allocated to project work in this unit. Members of a group will receive the same mark, unless the group collectively sign a letter, specifying how the workload was distributed. We very much encourage the use of the LMS discussion forum and class time for discussions of ideas. On the other hand, soliciting of help from outside of our small community will be considered cheating and will lead to disciplinary action.

The marking sheet for Stage 3 will be made available on the LMS. Marks will be awarded on the basis of correctness (some 80%) and programming structure, style, readability, commenting and layout (some 20%). Of the correctness marks, most will be directed towards code generation, with scanning, parsing, semantic analysis and symbol table handling counting for less.

## Appendix: Code format rules

Your Haskell programs should adhere with the following simple formatting rules:

- Each file should identify the team that produced it.

- Every non-trivial Haskell function should contain a comment at the beginning explaining its purpose and behaviour.

- Variable and function names must be meaningful.

- Significant blocks of code must be commented. However, not every statement in a program needs to be commented. Just as you can write too few comments, it is possible to write too many comments.

- Code should be laid out not only to satisfy Haskell's layout rules, but also to look neat. Program blocks appearing in if-expressions, let clauses, etc., must be indented consistently. They can be indented using tabs or spaces, and can be indented 2, 4, or 8 spaces, as long as it is done consistently. (Tabs can sometimes be rendered differently by different printers/viewers; spaces are more likely to preserve the intended layout.)

- Each program line should contain no more than 80 characters.

<div align="right">

Harald Søndergaard
1 May 2019

</div>